

Overview

- Previously we have seen how to learn a linear regression model using distributed closed-form solution.
- This distributed solution runs into minor issues when we try to scale.

Limitations of the Closed-Form OLS: Solving for the θ s

- Problem is equivalent to inverting $X'X$ matrix.
 - Inverse does not exist if matrix is not of full rank.
 - E.g., if one column is a linear combination of another (collinearity).
 - Note that $X'X$ is closely related to the covariance of the x data.
 - So we are in trouble if two or more variables are perfectly correlated.
 - Numerical problems can also occur if variables are almost collinear.
- Equivalent to solving a system of p linear equations.
 - Many good numerical methods for doing this, e.g., Gaussian elimination, LU decomposition, etc.
 - These are numerically more stable than direct inversion.
- Matrix inversion is not easily parallelized.

Alternative Ways of Solving OLS

- Closed-form solution
 - In this method, we will minimize RSS by explicitly taking its derivatives with respect to the β_j 's (sometimes written as w , the weight vector) and setting them to zero.
 - Do this via calculus with matrices.


$$\theta^* = (X^T X)^{-1} X^T y \rightarrow$$

- Gradient descent gives another way of minimizing RSS.
- Bayesian approach.
- Quadratic programming.
- Others.

Linear Regression via Gradient Descent

Given: Minimize $f(x)$. $J(w, X, y) = \text{Minimize } \sum_{i=1}^m (w^T X_i - y_i)^2$

Step 1: Find the zeros of the gradient function $f'(x)$ using bisection method, Newton-Raphson, or gradient descent.



RSS = Variance of ϵ

$$0 = \frac{\partial \sum \hat{\epsilon}_i^2}{\partial W} = \frac{\partial \left(\sum_{j=1}^n (X_j W - y_i)^2 \right)}{\partial W}$$

Gradient vector of partial derivatives

$$\nabla J(W) = \left(\sum_{j=1}^n (X_j W - y_i) X_j \right)$$

Pull model closer to examples with biggest residual.

$$W^{t+1} = W^t - \alpha^i \left(\sum_{j=1}^n (X_j W^t - y_i) X_j \right)$$

Gradient descent

For another derivation see: <http://cs229.stanford.edu/notes/cs229-notes1.pdf>.

OLS Using Gradient Descent

$J_q(W, X, 1, m) = \text{Minimize } \sum_{i=1}^m (W^T X_i - y_i)^2$ OLS objective function with decision variables W

- Initialize W = vector of zeros.
- Repeat until convergence.

$W_{t+1} = W_t - \alpha \frac{1}{n} \sum_{j=1}^n (X_j W_t - y_j) X_j$ OLS batch update rule

- End repeat.

True gradient is approximated by the gradient of the cost function only evaluated at all examples; adjust parameters proportional to this approximate gradient.

Intuitively, drag weight vector closer to the incorrectly predicted examples.

Batch vs. Stochastic


- That was batch-based descent.
- Stochastic gradient descent vs. batch.

OLS Using Gradient Descent

Stochastic gradient descent

$$\nabla J_{w_j}(W^t)$$

Partial derivative WRT to
variable w_j of error function
 $J(W)$ at point w^t



Stochastic update (after each example)

Let $W = (0, 0, \dots)$

Repeat

For j in $0 \dots n$ #each variable

For i in $1 \dots m$ #each example

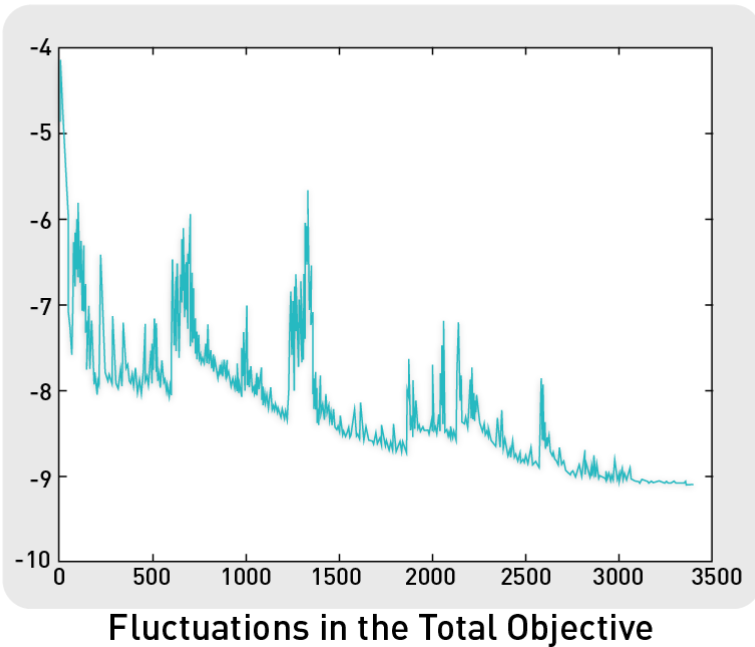
$$W^{t+1} = W^t - \alpha^i \left(\sum_{j=1}^n (X_j W^t - y_i) X_j \right)$$

Until convergence (i.e., no big changes in W or error)

Stochastic Gradient Descent vs. Batch

- Stochastic gradient descent can start making progress right away and continues to make progress with each example it looks at.
- Often, stochastic gradient descent gets W "close" to the minimum much faster than batch gradient descent.
 - Note, however, that it may never "converge" to the minimum, and the parameters W will keep oscillating around the minimum of $J(W)$; but in practice most of the values near the minimum will be reasonably good approximations to the true minimum.
- For these reasons, particularly when the training set is large, stochastic gradient descent is often preferred over batch gradient descent.

Stochastic Updates -> Fluctuations in the Total Objective



Two Approaches

- Two gradient-descent approaches to linear regression
 - Batch gradient descent
 - Stochastic gradient descent based
- Are both amenable to parallelization? And if so, which do you expect to perform better?

Stochastic Gradient Descent: Not Easy in MapReduce

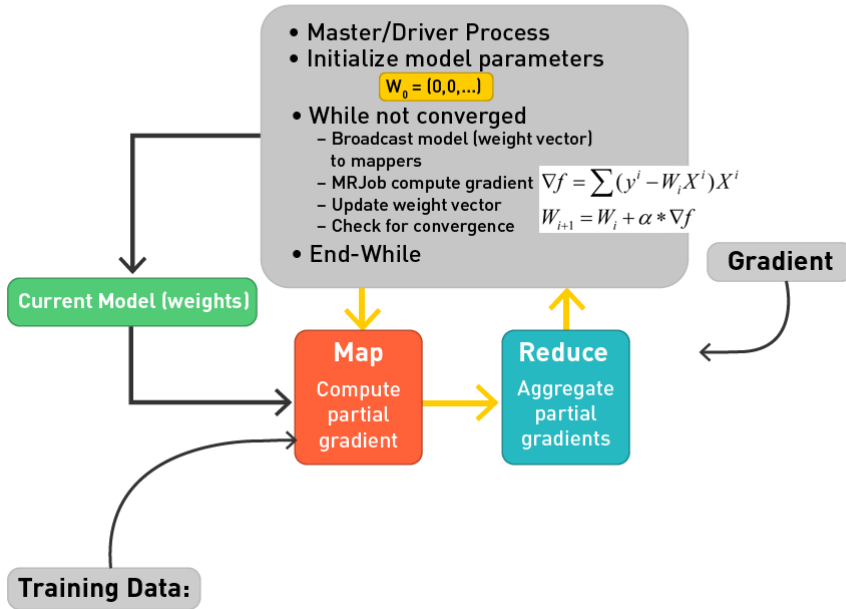
- Stochastic gradient descent requires the weight vector to be updated after each example (or minibatch of examples).
- "Some implementations of machine-learning algorithms, such as ICA, are commonly done with stochastic gradient ascent, which poses a challenge to parallelization. The problem is that in every step of gradient ascent, the algorithm updates a common set of parameters."

Some implementations of machine-learning algorithms, such as ICA, are commonly done with stochastic gradient ascent, which poses a challenge to parallelization. The problem is that in every step of gradient ascent, the algorithm updates a common set of parameters (e.g., the unmixing W matrix in ICA). When one gradient ascent step (involving one training sample) is updating W , it has to lock down this matrix, read it, compute the gradient, update W , and finally release the lock. This "lock-release" block creates a bottleneck for parallelization; thus, instead of stochastic gradient ascent, our algorithms above were implemented using batch gradient ascent.

OLS via Distributed Gradient Descent

- Master/Driver process.
- Initialize model parameters $W_0 = (0, 0, \dots)$.
- While not converged:
 - Broadcast model (weight vector) to mappers.
 - MRJob compute gradient $\nabla f = \sum (y_i - W_i X_i) X_i$.
 - Update weight vector $W_{i+1} = W_i + \alpha \nabla f$.
 - Check for convergence.
- End-While.

MapReduce Implementation



OLS via Distributed Gradient Descent: Mapper and Reducer

- Master/Driver process.
- Initialize model parameters, $W =$ vector of zeros: $W_0 = (0, 0, \dots)$.
- While not converged:
 - Broadcast model (i.e., weight vector) to the worker nodes.
 - Mapper (MANY mappers).
 - Compute partial gradient for each training example.
 - Combine in memory $\nabla f = \sum (y_i - W_i X_i) X_i$.
 - Finally yield the partial gradient.
 - Reducer (single Reducer).
 - Aggregate partial gradients.
 - Yield full gradient $\nabla f = \sum (y_i - W_i X_i) X_i$.
 - Update weight vector $W_{i+1} = W_i + \alpha * \nabla f$.
 - Check for convergence.
- End-While.

Divide and Conquer

- Can we use a combiner? Yes!
 - In-memory combiner.
 - So use mapper final.
- How many reducers?
 - We can have one (should be sufficient).
 - Or many (in the case of many, the master has to aggregate the partial reducer aggregates).

Linear Regression via Gradient Descent in MapReduce

