

13.4 A Supervised Random Walk

DATASCI W261

Machine Learning at Scale

Link Prediction Problem

Link Prediction Problem

- A fundamental problem in networks

Link Prediction Problem

- A fundamental problem in networks
- Here, given a snapshot of a network, need to infer:

Link Prediction Problem

- A fundamental problem in networks
- Here, given a snapshot of a network, need to infer:
 - Which interactions among members is likely

Link Prediction Problem

- A fundamental problem in networks
- Here, given a snapshot of a network, need to infer:
 - Which interactions among members is likely
 - Which existing interactions we are missing

Link Prediction Problem

- A fundamental problem in networks
- Here, given a snapshot of a network, need to infer:
 - Which interactions among members is likely
 - Which existing interactions we are missing
- Challenge: Combining information from network with rich node and edge attribute data

facebook Search Home Profile

Find friends from different parts of your life
Use the checkboxes below to discover people you know from your hometown, school, employer and more.

Hometown
☐ Indianapolis, Indiana
Enter another city

Current City
☐ Indianapolis, Indiana
Enter another city




High School
☐ North Central High School
Enter another high school

Mutual Friend
Enter a name

College or University
☐ Martin University
Enter another college

Employer
☐ ARIES GRAPHIC DESIGN
Enter another employer

Growing

 Judy Pyles 36 mutual friends Add Friend	 Rocky Campbell 41 mutual friends Add Friend	 Laura White 12 mutual friends Add Friend	 King Ro Conley 59 mutual friends Add Friend	 Dillon Rhodes 43 mutual friends Add Friend	 Rhonda Landrum 54 mutual friends Add Friend
 David Corbett 90 mutual friends Add Friend	 Eric Bettis 15 mutual friends Add Friend	 Eric Hughes 110 mutual friends Add Friend	 Markl Ann 26 mutual friends Add Friend	 Michael Pugh 21 mutual friends Add Friend	 Lisa Williams 22 mutual friends Add Friend
 LouieBaur Digg 39 mutual friends Add Friend	 LaTonya Mayberry Bynum 51 mutual friends Add Friend	 Durece Johnson 2 mutual friends Add Friend	 Kendale Adams 64 mutual friends Add Friend	 Bruce T. Caldwell 143 mutual friends Add Friend	 Angela Blackwell Miller 61 mutual friends Add Friend
 Landon Montel	 Kevin Brown	 Stanley F. Henry	 Saundria Mcrackin	 Ebonye X-Endsley	 Anita Hawkins



How do we combine network structure with node and edge features?

Supervised Link Prediction

- Combination of PageRank with supervised learning

Supervised Link Prediction

- Combination of PageRank with supervised learning
 - **PageRank:** Can capture importance of nodes based on network structure

Supervised Link Prediction

- Combination of PageRank with supervised learning
 - **PageRank:** Can capture importance of nodes based on network structure
 - **Supervised learning:** Uses node and edge features to adjust PageRank

Supervised Link Prediction

- Combination of PageRank with supervised learning
 - **PageRank:** Can capture importance of nodes based on network structure
 - **Supervised learning:** Uses node and edge features to adjust PageRank
- Idea: To "guide" random walk using supervised learning

Supervised Random Walk: Graph + Nodes + Edge Features

- Algorithm developed based on supervised random walks

Supervised Random Walk: Graph + Nodes + Edge Features

- Algorithm developed based on supervised random walks
 - Naturally combines information from the network with node and edge attributes

Supervised Random Walk: Graph + Nodes + Edge Features

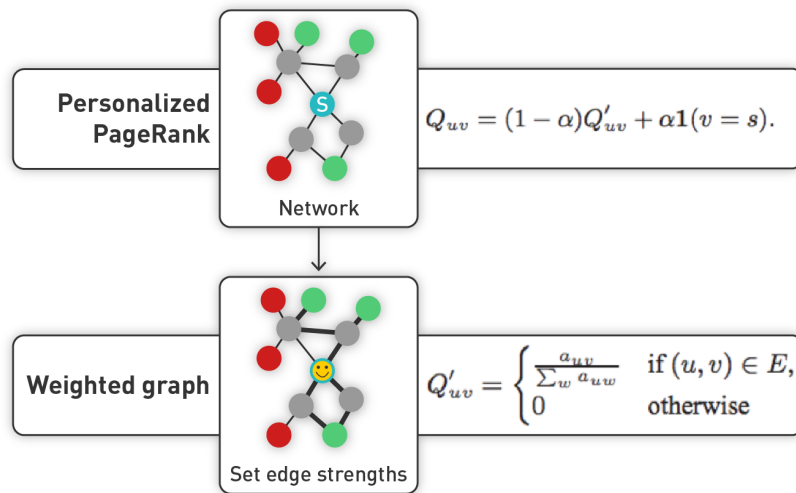
- Algorithm developed based on supervised random walks
 - Naturally combines information from the network with node and edge attributes
- Achieved through using attributes to guide a random walk on the graph

Supervised Random Walk: Graph + Nodes + Edge Features


- Algorithm developed based on supervised random walks
 - Naturally combines information from the network with node and edge attributes
- Achieved through using attributes to guide a random walk on the graph
- Problem formulated as part supervised machine learning and part random walking with restarts

Friend Graphs: Random Walk With Restarts

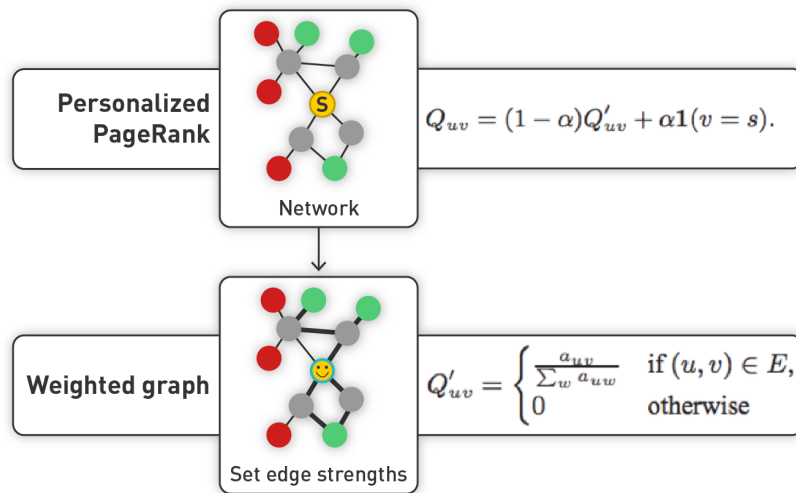
Friend Graphs: Random Walk With Restarts



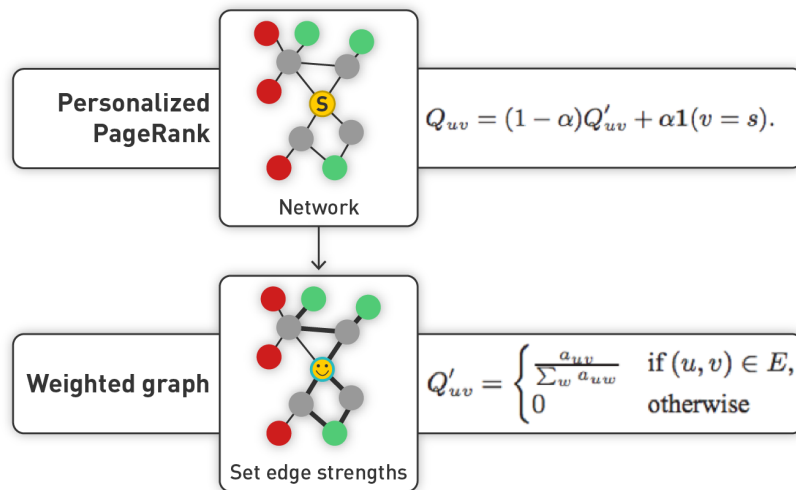
Friend Graphs: Random Walk With Restarts



Friend Graphs: Random Walk With Restarts




Friend Graphs: Random Walk With Restarts



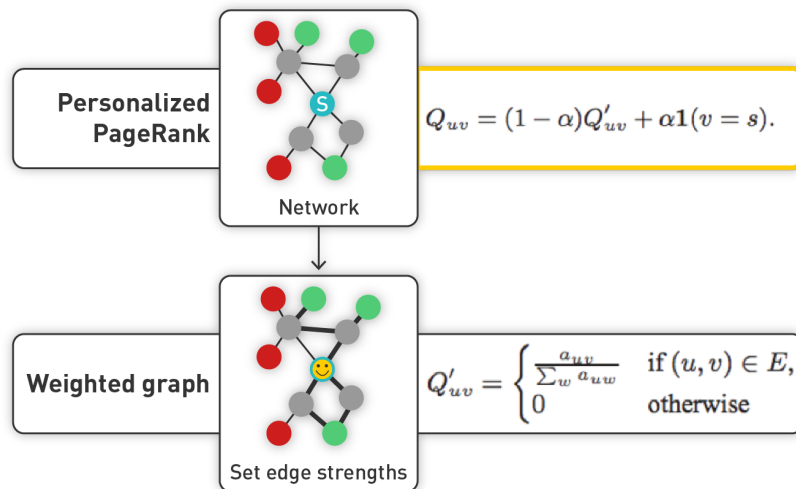
- Set teleportation factor to always teleport back to user.

Friend Graphs: Random Walk With Restarts



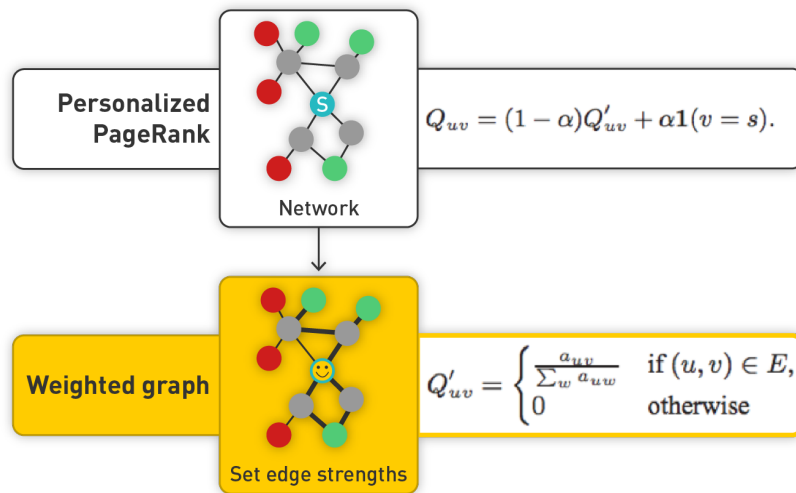
- Set teleportation factor to always teleport back to user.

Friend Graphs: Random Walk With Restarts



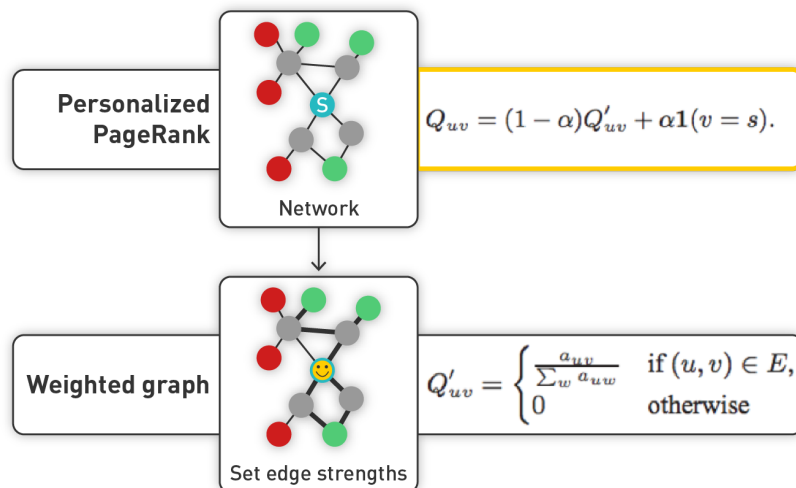
- Set teleportation factor to always teleport back to user.

Friend Graphs: Random Walk With Restarts




- Set teleportation factor to always teleport back to user.
- Reweight edges using supervised machine learning approach.

Friend Graphs: Random Walk With Restarts




- Set teleportation factor to always teleport back to user.
- Reweight edges using supervised machine learning approach.

Friend Graphs: Random Walk With Restarts




- Set teleportation factor to always teleport back to user.
- Reweight edges using supervised machine learning approach.

Friend Graphs: Random Walk With Restarts



- Set teleportation factor to always teleport back to user.
- Reweight edges using supervised machine learning approach.
- To distribute probability mass, use weighted sum distribution.

Friend Graphs: Random Walk With Restarts



- Set teleportation factor to always teleport back to user.
- Reweight edges using supervised machine learning approach.
- To distribute probability mass, use weighted sum distribution.

Supervised Learning

Supervised Learning


- Goal: Given a user s , recommend friends

Supervised Learning

- Goal: Given a user s , recommend friends

Positive: Nodes to which s links to in the future

Negative: Nodes to which s does not link during this future time period




Supervised Learning

- Goal: Given a user s , recommend friends

Positive: Nodes to which s links to in the future

Negative: Nodes to which s does not link during this future time period



Social Network Paper

By Lars Backstrom and Jure Leskovic ([link](#))

Social Network Paper

By Lars Backstrom and Jure Leskovic ([link](#))

- Sample: 200 users

Social Network Paper

By Lars Backstrom and Jure Leskovic ([link](#))

- Sample: 200 users
 - With a couple hundred friends

Social Network Paper

By Lars Backstrom and Jure Leskovic ([link](#))

- Sample: 200 users
 - With a couple hundred friends
 - Active at connecting

Social Network Paper

By Lars Backstrom and Jure Leskovic ([link](#))

- Sample: 200 users
 - With a couple hundred friends
 - Active at connecting
- Looked at time sequence of links

Social Network Paper

By Lars Backstrom and Jure Leskovic ([link](#))

- Sample: 200 users
 - With a couple hundred friends
 - Active at connecting
- Looked at time sequence of links
 - Positive and negative links selected

Social Network Paper

By Lars Backstrom and Jure Leskovic ([link](#))

- Sample: 200 users
 - With a couple hundred friends
 - Active at connecting
- Looked at time sequence of links
 - Positive and negative links selected
- Further partitioned positive and negative examples by time

Social Network Paper

By Lars Backstrom and Jure Leskovic ([link](#))

- Sample: 200 users
 - With a couple hundred friends
 - Active at connecting
- Looked at time sequence of links
 - Positive and negative links selected
- Further partitioned positive and negative examples by time
- Pooled all positive and negative examples for users, put in training and test set

Facebook Study Features

- Seven features generated around each edge (i, j)

Facebook Study Features

- Seven features generated around each edge (i, j)
 - Edge age: $T - t - \beta$, where T is time cutoff November 1, and t is edge creation time. Three features where $\beta = 0.10.30.5$.

Facebook Study Features

- Seven features generated around each edge (i, j)
 - Edge age: $T - t - \beta$, where T is time cutoff November 1, and t is edge creation time. Three features where $\beta = 0.10.30.5$.
 - Edge initiator: Individual making friend request encoded as +1 or -1

Facebook Study Features

- Seven features generated around each edge (i, j)
 - Edge age: $T - t - \beta$, where T is time cutoff November 1, and t is edge creation time. Three features where $\beta = 0.10.30.5$.
 - Edge initiator: Individual making friend request encoded as +1 or -1
 - Communication and observation features: Probability within a one-week period

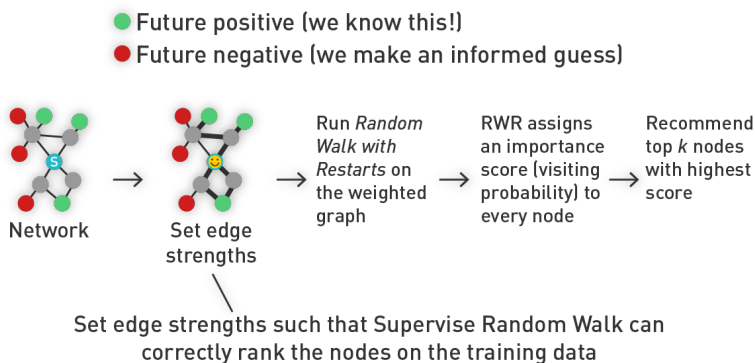
Facebook Study Features

- Seven features generated around each edge (i, j)
 - Edge age: $T - t - \beta$, where T is time cutoff November 1, and t is edge creation time. Three features where $\beta = 0.10.30.5$.
 - Edge initiator: Individual making friend request encoded as +1 or -1
 - Communication and observation features: Probability within a one-week period
 - Common friends: Between j and s

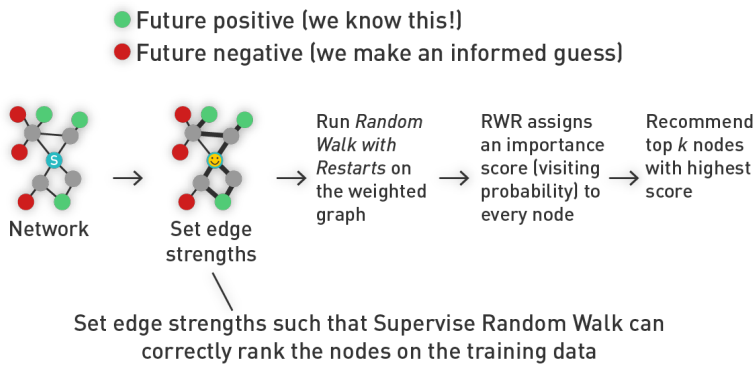
Facebook Study Features

- Seven features generated around each edge (i, j)
 - Edge age: $T - t - \beta$, where T is time cutoff November 1, and t is edge creation time. Three features where $\beta = 0.10.30.5$.
 - Edge initiator: Individual making friend request encoded as +1 or -1
 - Communication and observation features: Probability within a one-week period
 - Common friends: Between j and s
- All features rescaled to have mean 0 and standard deviation 1; also constant feature of value 1

Hybrid Model

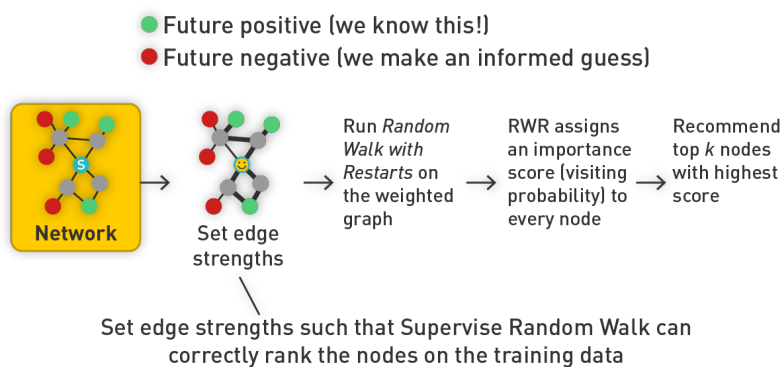


Hybrid Model



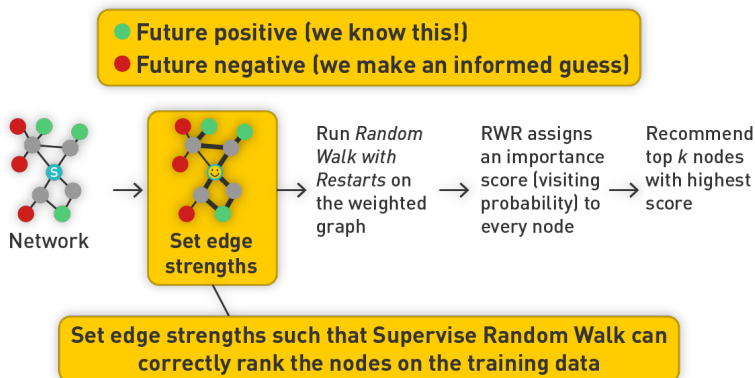
- Random walk with restarts algorithm combined with supervised learning

Hybrid Model



- Random walk with restarts algorithm combined with supervised learning
- Extremely complex, with billions of nodes

Hybrid Model



- Random walk with restarts algorithm combined with supervised learning
- Extremely complex, with billions of nodes
- Individual with positive and negative examples to reweight links and social graph

Supervised Random Walk: Process

Supervised Random Walk: Process

- Assume N individuals ($N = 100$) → 100 social network graphs.

Supervised Random Walk: Process

- Assume N individuals ($N = 100$) → 100 social network graphs.
- Weight edges applying weight vector W (seven dimensions) to all existing edges.

Supervised Random Walk: Process

- Assume N individuals ($N = 100$) \rightarrow 100 social network graphs.
- Weight edges applying weight vector W (seven dimensions) to all existing edges.
- Run random walk with restarts (RWR) algorithm (PPR with one node of interest).

$$Q_{uv} = (1 - \alpha)Q'_{uv} + \alpha \mathbf{1}(v = s)$$

Supervised Random Walk: Process

- Assume N individuals ($N = 100$) \rightarrow 100 social network graphs.
- Weight edges applying weight vector W (seven dimensions) to all existing edges.
- Run random walk with restarts (RWR) algorithm (PPR with one node of interest).

$$Q_{uv} = (1 - \alpha)Q'_{uv} + \alpha \mathbf{1}(v = s)$$

- Generate steady-state distribution using RWR.

Supervised Random Walk: Process

- Assume N individuals ($N = 100$) \rightarrow 100 social network graphs.
- Weight edges applying weight vector W (seven dimensions) to all existing edges.
- Run random walk with restarts (RWR) algorithm (PPR with one node of interest).

$$Q_{uv} = (1 - \alpha)Q'_{uv} + \alpha \mathbf{1}(v = s)$$

- Generate steady-state distribution using RWR.
- Teleport to s (our node of interest) at every opportunity.

Supervised Random Walk: Process (cont.)

- Steady-state distribution, where each node is scored

Supervised Random Walk: Process (cont.)

- Steady-state distribution, where each node is scored
- Leads to a list of nodes and corresponding scores

Supervised Random Walk: Process (cont.)

- Steady-state distribution, where each node is scored
- Leads to a list of nodes and corresponding scores
- Take labels and scores and use in training:

Supervised Random Walk: Process (cont.)

- Steady-state distribution, where each node is scored
- Leads to a list of nodes and corresponding scores
- Take labels and scores and use in training:
 - If node is positive, then positive example

Supervised Random Walk: Process (cont.)

- Steady-state distribution, where each node is scored
- Leads to a list of nodes and corresponding scores
- Take labels and scores and use in training:
 - If node is positive, then positive example
 - If node is negative, then negative example

Supervised Random Walk: Process (cont.)

- Steady-state distribution, where each node is scored
- Leads to a list of nodes and corresponding scores
- Take labels and scores and use in training:
 - If node is positive, then positive example
 - If node is negative, then negative example

$$\arg \min_{\theta} \sum_{p \in P} \sum_{n \in N} \delta(r_p < r_n) + \lambda \|\theta\|^2$$

Positive nodes \nearrow $p \in P$ Negative nodes \nwarrow $n \in N$
 r_x ... score of node x on a weighted graph with edge weights $f_{\theta}(x, y)$
 Penalty for violating constraint $r_p > r_n$

- Ideal: Score for positive examples > Score for negative examples ($r_p > r_n$)

Supervised Random Walk: Process (cont.)

- Steady-state distribution, where each node is scored
- Leads to a list of nodes and corresponding scores
- Take labels and scores and use in training:
 - If node is positive, then positive example
 - If node is negative, then negative example

$$\arg \min_{\theta} \sum_{p \in P} \sum_{n \in N} \delta(r_p < r_n) + \lambda \|\theta\|^2$$

Positive nodes $p \in P$ Negative nodes $n \in N$

r_x ... score of node x on a weighted graph with edge weights $f_{\theta}(x, y)$

Penalty for violating constraint $r_p > r_n$

- Ideal: Score for positive examples > Score for negative examples ($r_p > r_n$)
- Error: Score for negative examples > Score for positive examples ($r_p < r_n$)

Supervised Portion of Algorithm

Optimization function:

$$\min_w F(w) = \|w\|^2 + \lambda \sum_{d \in D, l \in L} h(p_l - p_d)$$

Supervised Portion of Algorithm

Optimization function:

$$\min_w F(w) = \|w\|^2 + \lambda \sum_{d \in D, l \in L} h(p_l - p_d)$$

Supervised Portion of Algorithm

Optimization function:

$$\min_w F(w) = ||w||^2 + \lambda \sum_{d \in D, l \in L} h(p_l - p_d)$$

- λ trades off between complexity (i.e., norm of w) for the fit of model (i.e., how much constraints can be violated).

Supervised Portion of Algorithm

Optimization function:

$$\min_w F(w) = ||w||^2 + \lambda \sum_{d \in D, l \in L} h(p_l - p_d)$$

- λ trades off between complexity (i.e., norm of w) for the fit of model (i.e., how much constraints can be violated).
- Apply gradient descent to objective function.

Supervised Portion of Algorithm

Optimization function:

$$\min_w F(w) = ||w||^2 + \lambda \sum_{d \in D, l \in L} h(p_l - p_d)$$

- λ trades off between complexity (i.e., norm of w) for the fit of model (i.e., how much constraints can be violated).
- Apply gradient descent to objective function.
- Learn a set of weights that can be applied to each edge in our graph.

Supervised Portion of Algorithm

Optimization function:

$$\min_w F(w) = ||w||^2 + \lambda \sum_{d \in D, l \in L} h(p_l - p_d)$$

- λ trades off between complexity (i.e., norm of w) for the fit of model (i.e., how much constraints can be violated).
- Apply gradient descent to objective function.
- Learn a set of weights that can be applied to each edge in our graph.
- Begin another iteration over newly rerated graph.

Supervised Portion of Algorithm

Optimization function:

$$\min_w F(w) = ||w||^2 + \lambda \sum_{d \in D, l \in L} h(p_l - p_d)$$

- λ trades off between complexity (i.e., norm of w) for the fit of model (i.e., how much constraints can be violated).
- Apply gradient descent to objective function.
- Learn a set of weights that can be applied to each edge in our graph.
- Begin another iteration over newly rerated graph.
- Generate new steady-state distribution; revisit problem. Readjust weights if needed.

Supervised Portion of Algorithm

Optimization function:

$$\min_w F(w) = ||w||^2 + \lambda \sum_{d \in D, l \in L} h(p_l - p_d)$$

- λ trades off between complexity (i.e., norm of w) for the fit of model (i.e., how much constraints can be violated).
- Apply gradient descent to objective function.
- Learn a set of weights that can be applied to each edge in our graph.
- Begin another iteration over newly rerated graph.
- Generate new steady-state distribution; revisit problem. Readjust weights if needed.
- Repeat process many times.

Implementation: Two Parts

Implementation: Two Parts

$$Q'_{uv} = \begin{cases} \frac{a_{uv}}{\sum_w a_{uw}} & \text{if } (u, v) \in E, \\ 0 & \text{otherwise} \end{cases}$$

- Supervised personalized random walk can be done on MapReduce

Implementation: Two Parts

$$Q'_{uv} = \begin{cases} \frac{a_{uv}}{\sum_w a_{uw}} & \text{if } (u, v) \in E, \\ 0 & \text{otherwise} \end{cases}$$

- Supervised personalized random walk can be done on MapReduce
 - Or using specialized libraries for graphs

Implementation: Two Parts

$$Q'_{uv} = \begin{cases} \frac{a_{uv}}{\sum_w a_{uw}} & \text{if } (u, v) \in E, \\ 0 & \text{otherwise} \end{cases}$$

- Supervised personalized random walk can be done on MapReduce
 - Or using specialized libraries for graphs
- Use gradient descent algorithm on a single server

$$\arg \min_{\theta} \sum_{p \in P} \sum_{n \in N} \delta(r_p < r_n) + \lambda \|\theta\|^2$$

Positive nodes $p \in P$ (green arrow)
 Negative nodes $n \in N$ (red arrow)
 $r_x \dots$ score of node x on a weighted graph with edge weights $f_{\theta}(x, y)$ (blue arrow)
 Penalty for violating constraint $r_p > r_n$ (blue arrow)